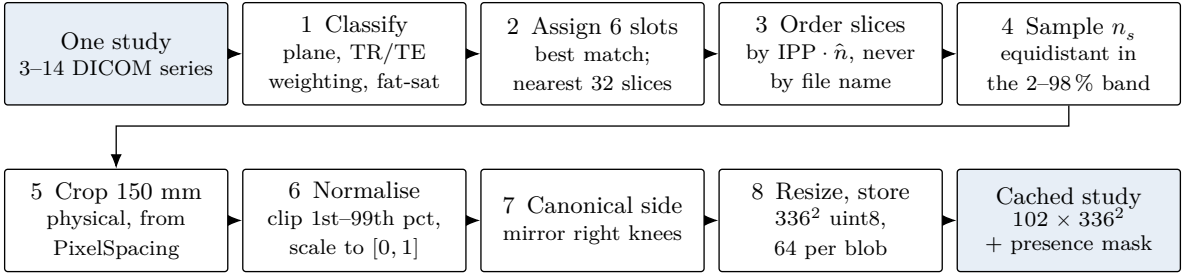

Slot budgets: SAG fluid FS 24, COR fluid FS 24, AX fluid FS 24, SAG fluid no FS 14, COR T1 8, SAG T1 8.
All 4,407 training studies: ≈ 51 GB, built once on Kaggle CPU kernels.
Hidden-test studies go through the same function at inference.